

Introduction to the tidyverse: data manipulation with `dplyr`

Christopher Skovron
Northwestern University

July 16, 2018

dplyr: data manipulation

dplyr is the main workhorse of the tidyverse

This lesson draws on [Chapter 5 of R for Data Science](#). For a condensed version, check out the [Introduction to dplyr vignette](#).

Example data: nycflights13

```
#install.packages(c('nycflights13', 'tidyverse'))  
library(nycflights13)  
library(tidyverse)
```

```
## — Attaching packages ————— tidyverse 1.2.1 —
```

```
## ✓ ggplot2 2.2.1      ✓ purrr 0.2.5  
## ✓ tibble 1.4.2       ✓ dplyr 0.7.5  
## ✓ tidyr 0.8.1        ✓ stringr 1.3.1  
## ✓ readr 1.1.1       ✓ forcats 0.3.0
```

```
## — Conflicts ————— tidyverse_conflicts() —  
## ✗ dplyr::filter() masks stats::filter()  
## ✗ dplyr::lag()     masks stats::lag()
```

Note conflicts when you load the tidyverse

Take careful note of the conflicts message that's printed when you load the tidyverse. It tells you that dplyr overwrites some functions in base R. If you want to use the base version of these functions after loading dplyr, you'll need to use their full names: `stats::filter()` and `stats::lag()`.

nycflights13

Example data: `nycflights13::flights`. This data frame contains all 336,776 flights that departed from New York City in 2013.

```
flights
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515             2     830
## 2  2013     1     1     533           529             4     850
## 3  2013     1     1     542           540             2     923
## 4  2013     1     1     544           545            -1    1004
## 5  2013     1     1     554           600            -6     812
## 6  2013     1     1     554           558            -4     740
## 7  2013     1     1     555           600            -5     913
## 8  2013     1     1     557           600            -3     709
## 9  2013     1     1     557           600            -3     838
## 10 2013     1     1     558           600            -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

It prints as a tibble

You might also have noticed the row of three (or four) letter abbreviations under the column names. These describe the type of each variable:

- `int` stands for integers.
- `dbl` stands for doubles, or real numbers.
- `chr` stands for character vectors, or strings.
- `dt` stands for date-times (a date + a time).

Other kinds of variables you may see in a tibble

- `lgl` stands for logical, vectors that contain only `TRUE` or `FALSE`.
- `fctr` stands for factors, which R uses to represent categorical variables with fixed possible values.
- `date` stands for dates.

dplyr verbs

- Pick observations by their values (`filter()`).
- Reorder the rows (`arrange()`).
- Pick variables by their names (`select()`).
- Create new variables with functions of existing variables (`mutate()`).
- Collapse many values down to a single summary (`summarise()`).

We will frequently use these on grouped data

These can all be used in conjunction with `group_by ()` which changes the scope of each function from operating on the entire dataset to operating on it group-by-group. These six functions provide the verbs for a language of data manipulation.

How dplyr verbs work

All verbs work similarly:

1. The first argument is a data frame.
2. The subsequent arguments describe what to do with the data frame, using the variable names (without quotes).
3. The result is a new data frame.

Together these properties make it easy to chain together multiple simple steps to achieve a complex result. Let's dive in and see how these verbs work.

Choose rows with `filter()`

For example, we can select all flights on January 1st with:

```
filter(flights, month == 1, day == 1)
```

```
## # A tibble: 842 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
## 3  2013     1     1     542           540           2     923
## 4  2013     1     1     544           545          -1    1004
## 5  2013     1     1     554           600          -6     812
## 6  2013     1     1     554           558          -4     740
## 7  2013     1     1     555           600          -5     913
## 8  2013     1     1     557           600          -3     709
## 9  2013     1     1     557           600          -3     838
## 10 2013     1     1     558           600          -2     753
## # ... with 832 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Choose rows with `filter()`

When you run that line of code, dplyr executes the filtering operation and returns a new data frame. dplyr functions never modify their inputs, so if you want to save the result, you'll need to use the assignment operator, `<-`:

```
jan1 <- filter(flights, month == 1, day == 1)
```

Comparisons

To use filtering effectively, you have to know how to select the observations that you want using the comparison operators. R provides the standard suite: `>`, `>=`, `<`, `<=`, `!=` (not equal), and `==` (equal).

Remember to use `=` instead of `==` when testing for equality. When this happens you'll get an informative error:

```
filter(flights, month = 1)
```

```
## Error: `month` (`month = 1`) must not be named, do you need `==`?
```

Issues with floating point numbers

There's another common problem you might encounter when using `==`: floating point numbers. These results might surprise you!

```
sqrt(2) ^ 2 == 2
```

```
## [1] FALSE
```

```
1/49 * 49 == 1
```

```
## [1] FALSE
```

Use `near ()` instead

Computers use finite precision arithmetic (they obviously can't store an infinite number of digits!) so remember that every number you see is an approximation. Instead of relying on `==`, use `near ()`:

```
near(sqrt(2) ^ 2, 2)
```

```
## [1] TRUE
```

```
near(1 / 49 * 49, 1)
```

```
## [1] TRUE
```


Logical operators

Multiple arguments to `filter()` are combined with “and”: every expression must be true in order for a row to be included in the output. For other types of combinations, you’ll need to use Boolean operators yourself: `&` is “and”, `|` is “or”, and `!` is “not”.

Find all flights that departed in November or December:

```
filter(flights, month == 11 | month == 12)
```

Another way to write it using %in%

```
nov_dec <- filter(flights, month %in% c(11, 12))
```

Sometimes you can simplify complicated subsetting by remembering De Morgan's law: $!(x \ \& \ y)$ is the same as $!x \mid !y$, and $!(x \mid y)$ is the same as $!x \ \& \ !y$. For example, if you wanted to find flights that weren't delayed (on arrival or departure) by more than two hours, you could use either of the following two filters:

```
filter(flights, !(arr_delay > 120 | dep_delay > 120))  
filter(flights, arr_delay <= 120, dep_delay <= 120)
```

Missing values

If you want to determine if a value is missing, use `is.na()`:

```
is.na(x)
```

`filter()` only includes rows where the condition is `TRUE`; it excludes both `FALSE` and `NA` values. If you want to preserve missing values, ask for them explicitly:

```
df <- tibble(x = c(1, NA, 3))  
filter(df, x > 1)
```

```
## # A tibble: 1 x 1  
##       x  
##   <dbl>  
## 1     3
```

```
filter(df, is.na(x) | x > 1)
```

```
## # A tibble: 2 x 1
##       x
##   <dbl>
## 1     NA
## 2      3
```

Arrange rows with `arrange()`

`arrange()` changes the order of rows order. It takes a data frame and a set of column names (or more complicated expressions) to order by. If you provide more than one column name, each additional column will be used to break ties in the values of preceding columns:

```
arrange(flights, year, month, day)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
## 3  2013     1     1     542           540           2     923
## 4  2013     1     1     544           545          -1    1004
## 5  2013     1     1     554           600          -6     812
## 6  2013     1     1     554           558          -4     740
## 7  2013     1     1     555           600          -5     913
## 8  2013     1     1     557           600          -3     709
## 9  2013     1     1     557           600          -3     838
## 10 2013     1     1     558           600          -2     753
```

```
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <dbl>,  
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,  
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,  
## #   minute <dbl>, time_hour <dtm>
```

Arrange rows with `arrange()`

Use `desc()` to re-order by a column in descending order:

```
arrange(flights, desc(arr_delay))
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     9     641           900         1301    1242
## 2  2013     6    15    1432          1935         1137    1607
## 3  2013     1    10    1121          1635         1126    1239
## 4  2013     9    20    1139          1845         1014    1457
## 5  2013     7    22     845          1600         1005    1044
## 6  2013     4    10    1100          1900          960    1342
## 7  2013     3    17    2321           810          911     135
## 8  2013     7    22    2257           759          898     121
## 9  2013    12     5     756          1700          896    1058
## 10 2013     5     3    1133          2055          878    1250
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <dbl>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```


Arrange rows with `arrange()`

Missing values are always sorted at the end:

```
df <- tibble(x = c(5, 2, NA))  
arrange(df, x)
```

```
## # A tibble: 3 x 1  
##       x  
##   <dbl>  
## 1     2  
## 2     5  
## 3    NA
```

```
arrange(df, desc(x))
```

```
## # A tibble: 3 x 1  
##       x  
##   <dbl>  
## 1     5  
## 2     2  
## 3    NA
```

Exercise

1. Sort `flights` to find the fastest flights.

Solution

```
flights %>%  
  arrange(air_time)
```

```
## # A tibble: 336,776 x 19  
##   year month   day dep_time sched_dep_time dep_delay arr_time  
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>  
## 1  2013     1    16    1355           1315         40    1442  
## 2  2013     4    13     537           527         10     622  
## 3  2013    12     6     922           851         31    1021  
## 4  2013     2     3    2153          2129         24    2247  
## 5  2013     2     5    1303          1315        -12    1342  
## 6  2013     2    12    2123          2130         -7    2211  
## 7  2013     3     2    1450          1500        -10    1547  
## 8  2013     3     8    2026          1935         51    2131  
## 9  2013     3    18    1456          1329         87    1533  
## 10 2013     3    19    2226          2145         41    2305  
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <dbl>,  
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,  
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,  
## #   minute <dbl>, time_hour <dtm>
```

Select columns with `select()`

`select()` chooses the columns you want out of the whole set of columns

```
# Select columns by name  
select(flights, year, month, day)
```

```
## # A tibble: 336,776 x 3  
##   year month   day  
##   <int> <int> <int>  
## 1  2013     1     1  
## 2  2013     1     1  
## 3  2013     1     1  
## 4  2013     1     1  
## 5  2013     1     1  
## 6  2013     1     1  
## 7  2013     1     1  
## 8  2013     1     1  
## 9  2013     1     1  
## 10 2013     1     1  
## # ... with 336,766 more rows
```

Select columns with `select()`

```
# Select all columns between year and day (inclusive)
select(flights, year:day)
```

```
## # A tibble: 336,776 x 3
##   year month   day
##   <int> <int> <int>
## 1  2013     1     1
## 2  2013     1     1
## 3  2013     1     1
## 4  2013     1     1
## 5  2013     1     1
## 6  2013     1     1
## 7  2013     1     1
## 8  2013     1     1
## 9  2013     1     1
## 10 2013     1     1
## # ... with 336,766 more rows
```

Select columns with `select()`

```
# Select all columns except those from year to day (inclusive)
select(flights, -(year:day))
```

```
## # A tibble: 336,776 x 16
##   dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
##   <int>         <int>         <dbl>   <int>         <int>         <dbl>
## 1      517           515           2      830           819           11
## 2      533           529           4      850           830           20
## 3      542           540           2      923           850           33
## 4      544           545          -1     1004          1022          -18
## 5      554           600          -6      812           837          -25
## 6      554           558          -4      740           728           12
## 7      555           600          -5      913           854           19
## 8      557           600          -3      709           723          -14
## 9      557           600          -3      838           846           -8
## 10     558           600          -2      753           745            8
## # ... with 336,766 more rows, and 10 more variables: carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Helpers for `select ()`

There are a number of helper functions you can use within `select ()`:

- `starts_with("abc")`: matches names that begin with “abc”.
- `ends_with("xyz")`: matches names that end with “xyz”.
- `contains("ijk")`: matches names that contain “ijk”.
- `matches("(.)\\1")`: selects variables that match a regular expression. This one matches any variables that contain repeated characters. You’ll learn more about regular expressions in [strings].
- `num_range("x", 1:3)` matches `x1`, `x2` and `x3`.

Use `rename ()` to rename variables

`select ()` can be used to rename variables, but it's rarely useful because it drops all of the variables not explicitly mentioned. Instead, use `rename ()`, which is a variant of `select ()` that keeps all the variables that aren't explicitly mentioned:

Use `rename ()` to rename variables

```
rename(flights, tail_num = tailnum)
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
## 3  2013     1     1     542           540           2     923
## 4  2013     1     1     544           545          -1    1004
## 5  2013     1     1     554           600          -6     812
## 6  2013     1     1     554           558          -4     740
## 7  2013     1     1     555           600          -5     913
## 8  2013     1     1     557           600          -3     709
## 9  2013     1     1     557           600          -3     838
## 10 2013     1     1     558           600          -2     753
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time <dbl>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tail_num <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Note that it's `new_name = old_name`. This always throws me off.

Add new variables with `mutate()`

Besides selecting sets of existing columns, it's often useful to add new columns that are functions of existing columns.

That's the job of `mutate()`.

Make a smaller version of the dataset so we can see our work

```
flights_sml <- select(flights,  
  year:day,  
  ends_with("delay"),  
  distance,  
  air_time  
)
```

Add new variables with `mutate()`

```
mutate(flights_sml,  
  gain = arr_delay - dep_delay,  
  speed = distance / air_time * 60  
)
```

```
## # A tibble: 336,776 x 9  
##   year month   day dep_delay arr_delay distance air_time  gain speed  
##   <int> <int> <int>   <dbl>   <dbl>   <dbl>   <dbl> <dbl> <dbl>  
## 1  2013     1     1         2       11    1400    227     9   370  
## 2  2013     1     1         4       20    1416    227    16   374  
## 3  2013     1     1         2       33    1089    160    31   408  
## 4  2013     1     1        -1      -18    1576    183   -17   517  
## 5  2013     1     1        -6      -25     762    116   -19   394  
## 6  2013     1     1        -4       12     719    150    16   288  
## 7  2013     1     1        -5       19    1065    158    24   404  
## 8  2013     1     1        -3      -14     229     53   -11   259  
## 9  2013     1     1        -3       -8     944    140    -5   405  
## 10 2013     1     1        -2        8     733    138    10   319  
## # ... with 336,766 more rows
```

Add new variables with `mutate()`

Note that you can refer to columns that you've just created:

```
mutate(flights_sml,  
  gain = arr_delay - dep_delay,  
  hours = air_time / 60,  
  gain_per_hour = gain / hours  
)
```

```
## # A tibble: 336,776 x 10  
##   year month   day dep_delay arr_delay distance air_time  gain hours  
##   <int> <int> <int>     <dbl>     <dbl>     <dbl>     <dbl> <dbl> <dbl>  
## 1  2013     1     1         2         11     1400     227      9 3.78  
## 2  2013     1     1         4         20     1416     227     16 3.78  
## 3  2013     1     1         2         33     1089     160     31 2.67  
## 4  2013     1     1        -1        -18     1576     183    -17 3.05  
## 5  2013     1     1        -6        -25      762     116    -19 1.93  
## 6  2013     1     1        -4         12      719     150     16 2.5  
## 7  2013     1     1        -5         19     1065     158     24 2.63  
## 8  2013     1     1        -3        -14      229      53    -11 0.88  
## 9  2013     1     1        -3         -8      944     140     -5 2.33  
## 10 2013     1     1        -2          8      733     138     10 2.3  
## # ... with 336,766 more rows, and 1 more variable: gain_per_hour <dbl>
```

Add new variables with `mutate()`

If you only want to keep the new variables, use `transmute()`:

```
transmute(flights,  
  gain = arr_delay - dep_delay,  
  hours = air_time / 60,  
  gain_per_hour = gain / hours  
)
```

```
## # A tibble: 336,776 x 3  
##       gain hours gain_per_hour  
##   <dbl> <dbl>         <dbl>  
## 1      9 3.78           2.38  
## 2     16 3.78           4.23  
## 3     31 2.67          11.6  
## 4    -17 3.05          -5.57  
## 5    -19 1.93          -9.83  
## 6     16 2.5            6.4  
## 7     24 2.63           9.11  
## 8    -11 0.883         -12.5  
## 9     -5 2.33          -2.14  
## 10    10 2.3            4.35  
## # ... with 336,766 more rows
```

There are many functions for creating new variables that you can use with `mutate()`. The key property is that the function must be vectorised: it must take a vector of values as input, return a vector with the same number of values as output. There's no way to list every possible function that you might use, but here's a selection of functions that are frequently useful:

- Logical comparisons, `<`, `<=`, `>`, `>=`, `!=`, which you learned about earlier. If you're doing a complex sequence of logical operations it's often a good idea to store the interim values in new variables so you can check that each step is working as expected.

Useful helpers for `mutate()`

- Ranking: there are a number of ranking functions, but you should start with `min_rank()`. It does the most usual type of ranking (e.g. 1st, 2nd, 2nd, 4th). The default gives smallest values the small ranks; use `desc(x)` to give the largest values the smallest ranks.

```
y <- c(1, 2, 2, NA, 3, 4)
min_rank(y)
```

```
## [1] 1 2 2 NA 4 5
```

```
min_rank(desc(y))
```

```
## [1] 5 3 3 NA 2 1
```


Useful helpers for `mutate()`

If `min_rank()` doesn't do what you need, look at the variants `row_number()`, `dense_rank()`, `percent_rank()`, `cume_dist()`, `ntile()`. See their help pages for more details.

```
```r
row_number(y)
```

```
[1] 1 2 3 NA 4 5
```

```r
dense_rank(y)
```

```
[1] 1 2 2 NA 3 4
```

```r
percent_rank(y)
```

# Grouped summaries with `summarise()`

The last key verb is `summarise()`. It collapses a data frame to a single row:

```
summarise(flights, delay = mean(dep_delay, na.rm = TRUE))
```

```
A tibble: 1 x 1
delay
<dbl>
1 12.6
```

# Grouped summaries with `summarise()`

`summarise()` is not terribly useful unless we pair it with `group_by()`. This changes the unit of analysis from the complete dataset to individual groups. Then, when you use the dplyr verbs on a grouped data frame they'll be automatically applied "by group". For example, if we applied exactly the same code to a data frame grouped by date, we get the average delay per date:

```
by_day <- group_by(flights, year, month, day)
summarise(by_day, delay = mean(dep_delay, na.rm = TRUE))
```

```
A tibble: 365 x 4
Groups: year, month [?]
year month day delay
<int> <int> <int> <dbl>
1 2013 1 1 11.5
2 2013 1 2 13.9
3 2013 1 3 11.0
4 2013 1 4 8.95
```

```
4 2013 1 4 8.93
5 2013 1 5 5.73
6 2013 1 6 7.15
7 2013 1 7 5.42
8 2013 1 8 2.55
9 2013 1 9 2.28
10 2013 1 10 2.84
... with 355 more rows
```

## Missing values

You may have wondered about the `na.rm` argument we used above. What happens if we don't set it?

```
flights %>%
 group_by(year, month, day) %>%
 summarise(mean = mean(dep_delay))
```

```
A tibble: 365 x 4
Groups: year, month [?]
year month day mean
<int> <int> <int> <dbl>
1 2013 1 1 NA
2 2013 1 2 NA
3 2013 1 3 NA
4 2013 1 4 NA
5 2013 1 5 NA
6 2013 1 6 NA
7 2013 1 7 NA
```

```
/ 2013 1 / NA
8 2013 1 8 NA
9 2013 1 9 NA
10 2013 1 10 NA
... with 355 more rows
```

# Missing values

We get a lot of missing values! That's because aggregation functions obey the usual rule of missing values: if there's any missing value in the input, the output will be a missing value.

Fortunately, all aggregation functions have an `na.rm` argument which removes the missing values prior to computation:

```
flights %>%
 group_by(year, month, day) %>%
 summarise(mean = mean(dep_delay, na.rm = TRUE))
```

```
A tibble: 365 x 4
Groups: year, month [?]
year month day mean
<int> <int> <int> <dbl>
1 2013 1 1 11.5
2 2013 1 2 13.9
3 2013 1 3 11.0
4 2013 1 4 8.95
5 2013 1 5 5.73
```

```
6 2013 1 6 7.15
7 2013 1 7 5.42
8 2013 1 8 2.55
9 2013 1 9 2.28
10 2013 1 10 2.84
... with 355 more rows
```

# Missing values

In this case, where missing values represent cancelled flights, we could also tackle the problem by first removing the cancelled flights. We'll save this dataset so we can reuse in the next few examples.

```
not_cancelled <- flights %>%
 filter(!is.na(dep_delay), !is.na(arr_delay))

not_cancelled %>%
 group_by(year, month, day) %>%
 summarise(mean = mean(dep_delay))
```

```
A tibble: 365 x 4
Groups: year, month [?]
year month day mean
<int> <int> <int> <dbl>
1 2013 1 1 11.4
2 2013 1 2 13.7
3 2013 1 3 10.9
4 2013 1 4 8.97
5 2013 1 5 5.73
6 2013 1 6 7.15
```




```
7 2013 1 7 5.42
8 2013 1 8 2.56
9 2013 1 9 2.30
10 2013 1 10 2.84
... with 355 more rows
```

# Counts

Whenever you do any aggregation, it's always a good idea to include either a count (`n ( )`), or a count of non-missing values (`sum( !is.na(x) )`). That way you can check that you're not drawing conclusions based on very small amounts of data. For example, let's look at the planes (identified by their tail number) that have the highest average delays:

```
delays <- not_cancelled %>%
 group_by(tailnum) %>%
 summarise(
 delay = mean(arr_delay)
)

ggplot(data = delays, mapping = aes(x = delay)) +
 geom_freqpoly(binwidth = 10)
```



# Exercise

1. Find all flights that
  1. Had an arrival delay of two or more hours
  2. Flew to Houston (IAH or HOU)
  3. Were operated by United, American, or Delta
  4. Departed in summer (July, August, and September)

# Grouped summaries

# Counts

You've seen `n()`, which takes no arguments, and returns the size of the current group. To count the number of non-missing values, use `sum(!is.na(x))`. To count the number of distinct (unique) values, use `n_distinct(x)`.

```
Which destinations have the most carriers?
not_cancelled %>%
 group_by(dest) %>%
 summarise(carriers = n_distinct(carrier)) %>%
 arrange(desc(carriers))
```

```
A tibble: 104 x 2
dest carriers
<chr> <int>
1 ATL 7
2 BOS 7
3 CLT 7
4 ORD 7
5 TPA 7
6 AUS 6
7 DCA 6
8 DTW 6
```

```
9 IAD 6
10 MSP 6
... with 94 more rows
```

# Counts

Counts are so useful that dplyr provides a simple helper if all you want is a count:

```
not_cancelled %>%
 count(dest)
```

```
A tibble: 104 x 2
dest n
<chr> <int>
1 ABQ 254
2 ACK 264
3 ALB 418
4 ANC 8
5 ATL 16837
6 AUS 2411
7 AVL 261
8 BDL 412
9 BGR 358
10 BHM 269
... with 94 more rows
```



# Counts

You can optionally provide a weight variable. For example, you could use this to “count” (sum) the total number of miles a plane flew:

```
not_cancelled %>%
 count(tailnum, wt = distance)
```

```
A tibble: 4,037 x 2
tailnum n
<chr> <dbl>
1 D942DN 3418
2 N0EGMQ 239143
3 N10156 109664
4 N102UW 25722
5 N103US 24619
6 N104UW 24616
7 N10575 139903
8 N105UW 23618
9 N107US 21677
10 N108UW 32070
... with 4,027 more rows
```

# Counts

Counts and proportions of logical values: `sum(x > 10)`, `mean(y == 0)`. When used with numeric functions, TRUE is converted to 1 and FALSE to 0. This makes `sum()` and `mean()` very useful: `sum(x)` gives the number of TRUEs in `x`, and `mean(x)` gives the proportion.

```
How many flights left before 5am? (these usually indicate delayed
flights from the previous day)
not_cancelled %>%
 group_by(year, month, day) %>%
 summarise(n_early = sum(dep_time < 500))
```

```
A tibble: 365 x 4
Groups: year, month [?]
year month day n_early
<int> <int> <int> <int>
1 2013 1 1 0
2 2013 1 2 3
3 2013 1 3 4
4 2013 1 4 3
5 2013 1 5 2
```

```
5 2013 1 5 5
6 2013 1 6 2
7 2013 1 7 2
8 2013 1 8 1
9 2013 1 9 3
10 2013 1 10 3
... with 355 more rows
```

```
What proportion of flights are delayed by more than an hour?
not_cancelled %>%
 group_by(year, month, day) %>%
 summarise(hour_perc = mean(arr_delay > 60))
```

```
A tibble: 365 x 4
Groups: year, month [?]
year month day hour_perc
<int> <int> <int> <dbl>
1 2013 1 1 0.0722
2 2013 1 2 0.0851
3 2013 1 3 0.0567
4 2013 1 4 0.0396
5 2013 1 5 0.0349
6 2013 1 6 0.0470
7 2013 1 7 0.0333
8 2013 1 8 0.0213
9 2013 1 9 0.0202
10 2013 1 10 0.0183
... with 355 more rows
```

# Grouping by multiple variables

When you group by multiple variables, each summary peels off one level of the grouping. That makes it easy to progressively roll up a dataset:

```
daily <- group_by(flights, year, month, day)
(per_day <- summarise(daily, flights = n()))
```

```
A tibble: 365 x 4
Groups: year, month [?]
year month day flights
<int> <int> <int> <int>
1 2013 1 1 842
2 2013 1 2 943
3 2013 1 3 914
4 2013 1 4 915
5 2013 1 5 720
6 2013 1 6 832
7 2013 1 7 933
8 2013 1 8 899
9 2013 1 9 902
10 2013 1 10 932
... with 355 more rows
```

```
(per_month <- summarise(per_day, flights = sum(flights)))
```

```
A tibble: 12 x 3
Groups: year [?]
year month flights
<int> <int> <int>
1 2013 1 27004
2 2013 2 24951
3 2013 3 28834
4 2013 4 28330
5 2013 5 28796
6 2013 6 28243
7 2013 7 29425
8 2013 8 29327
9 2013 9 27574
10 2013 10 28889
11 2013 11 27268
12 2013 12 28135
```

```
(per_year <- summarise(per_month, flights = sum(flights)))
```

```
A tibble: 1 x 2
year flights
<int> <int>
1 2013 336776
```

Be careful when progressively rolling up summaries: it's OK for sums and counts, but you need to think about weighting means and variances, and it's not possible to do it exactly for rank-based statistics like the median. In other words, the sum of groupwise sums is the overall sum, but the median of groupwise medians is not the overall median.

# Ungrouping

If you need to remove grouping, and return to operations on ungrouped data, use `ungroup()`.

```
daily %>%
 ungroup() %>% # no longer grouped by date
 summarise(flights = n()) # all flights
```

```
A tibble: 1 x 1
flights
<int>
1 336776
```

# Grouped mutates (and filters)

Grouping is most useful in conjunction with `summarise()`,  
but you can also do convenient operations with `mutate()`  
and `filter()`:



# Find the worst members of each group:

```
```r
flights_sml %>%
  group_by(year, month, day) %>%
  filter(rank(desc(arr_delay)) < 10)
```

```
## # A tibble: 3,306 x 7
## # Groups:   year, month, day [365]
##   year month   day dep_delay arr_delay distance air_time
##   <int> <int> <int>     <dbl>     <dbl>     <dbl>     <dbl>
## 1  2013     1     1       853       851       184        41
## 2  2013     1     1       290       338      1134       213
## 3  2013     1     1       260       263       266        46
## 4  2013     1     1       157       174       213        60
## 5  2013     1     1       216       222       708       121
## 6  2013     1     1       255       250       589       115
## 7  2013     1     1       285       246     1085       146

```

Find all groups bigger than a threshold:

```
```r
popular_dests <- flights %>%
 group_by(dest) %>%
 filter(n() > 365)
popular_dests
```

```
A tibble: 332,577 x 19
Groups: dest [77]
year month day dep_time sched_dep_time dep_delay arr_time
<int> <int> <int> <int> <int> <dbl> <int>
1 2013 1 1 517 515 2 830
2 2013 1 1 533 529 4 850
3 2013 1 1 542 540 2 923
4 2013 1 1 544 545 -1 1004
5 2013 1 1 554 600 -6 812
6 2013 1 1 554 558 -4 740

```

# Standardise to compute per group metrics:

```
```r
popular_dests %>%
  filter(arr_delay > 0) %>%
  mutate(prop_delay = arr_delay / sum(arr_delay)) %>%
  select(year:day, dest, arr_delay, prop_delay)
```
```

```
```
```

```
## # A tibble: 131,106 x 6
## # Groups:   dest [77]
##   year month   day dest  arr_delay prop_delay
##   <int> <int> <int> <chr>    <dbl>    <dbl>
## 1  2013     1     1 IAH      11  0.000111
## 2  2013     1     1 IAH      20  0.000201
## 3  2013     1     1 MIA      33  0.000235
## 4  2013     1     1 ORD      12  0.0000424
## 5  2013     1     1 FLL      19  0.0000938
## 6  2013     1     1 ORD       8  0.0000283
```

Exercise

1. Come up with another approach that will give you the same output as `not_cancelled %>% count(dest)` and `not_cancelled %>% count(tailnum, wt = distance)` (without using `count()`).
2. Which carrier has the worst delays?

Advanced verbs: `summarize_all()`,
`summarize_at()`, `summarize_if()`

The `_all`, `_at` and `_if` suffixes work on `mutate`,
`summarize`, and `transmute`. (British English speakers, you
can use `summarise` too.)

- `_all` applies the operation to all columns
- `_at` applies the operation only to named columns
- `_if` applies the operation if some condition is true

summarize_all()

```
```r
iris %>%
 group_by(Species) %>%
 summarise_all(mean)
```

```
A tibble: 3 x 5
Species Sepal.Length Sepal.Width Petal.Length Petal.Width
<fct> <dbl> <dbl> <dbl> <dbl>
1 setosa 5.01 3.43 1.46 0.246
2 versicolor 5.94 2.77 4.26 1.33
3 virginica 6.59 2.97 5.55 2.03
```
```

summarize_at()

```
starwars %>%  
summarise_at(c("height", "mass"), mean, na.rm = TRUE)
```

```
## # A tibble: 1 x 2  
##   height  mass  
##   <dbl> <dbl>  
## 1   174.   97.3
```

mutate_at() with a helper to choose columns

```
iris %>%  
  mutate_at(vars(matches("Sepal")), log)
```

| ## | Sepal.Length | Sepal.Width | Petal.Length | Petal.Width | Species |
|-------|--------------|-------------|--------------|-------------|---------|
| ## 1 | 1.629241 | 1.2527630 | 1.4 | 0.2 | setosa |
| ## 2 | 1.589235 | 1.0986123 | 1.4 | 0.2 | setosa |
| ## 3 | 1.547563 | 1.1631508 | 1.3 | 0.2 | setosa |
| ## 4 | 1.526056 | 1.1314021 | 1.5 | 0.2 | setosa |
| ## 5 | 1.609438 | 1.2809338 | 1.4 | 0.2 | setosa |
| ## 6 | 1.686399 | 1.3609766 | 1.7 | 0.4 | setosa |
| ## 7 | 1.526056 | 1.2237754 | 1.4 | 0.3 | setosa |
| ## 8 | 1.609438 | 1.2237754 | 1.5 | 0.2 | setosa |
| ## 9 | 1.481605 | 1.0647107 | 1.4 | 0.2 | setosa |
| ## 10 | 1.589235 | 1.1314021 | 1.5 | 0.1 | setosa |
| ## 11 | 1.686399 | 1.3083328 | 1.5 | 0.2 | setosa |
| ## 12 | 1.568616 | 1.2237754 | 1.6 | 0.2 | setosa |
| ## 13 | 1.568616 | 1.0986123 | 1.4 | 0.1 | setosa |
| ## 14 | 1.458615 | 1.0986123 | 1.1 | 0.1 | setosa |
| ## 15 | 1.757858 | 1.3862944 | 1.2 | 0.2 | setosa |
| ## 16 | 1.740466 | 1.4816045 | 1.5 | 0.4 | setosa |
| ## 17 | 1.686399 | 1.3609766 | 1.3 | 0.4 | setosa |

summarise_if

```
# The _if() variants apply a predicate function (a function that  
# returns TRUE or FALSE) to determine the relevant subset of  
# columns. Here we apply mean() to the numeric columns:  
starwars %>% summarise_if(is.numeric, mean, na.rm = TRUE)
```

```
## # A tibble: 1 x 3  
##   height mass birth_year  
##   <dbl> <dbl>      <dbl>  
## 1   174.  97.3        87.6
```

```
#> # A tibble: 1 x 3  
#>   height mass birth_year  
#>   <dbl> <dbl>      <dbl>  
#> 1   174.  97.3        87.6
```

Multiple grouped transformations

```
# If you want to apply multiple transformations, use funs()  
by_species <- iris %>% group_by(Species)  
  
by_species %>% summarise_all(funs(min, max))
```

```
## # A tibble: 3 x 9  
##   Species      Sepal.Length_min Sepal.Width_min Petal.Length_min  
##   <fct>          <dbl>          <dbl>          <dbl>  
## 1 setosa         4.3            2.3            1  
## 2 versicolor    4.9            2             3  
## 3 virginica      4.9            2.2           4.5  
## # ... with 5 more variables: Petal.Width_min <dbl>,  
## #   Sepal.Length_max <dbl>, Sepal.Width_max <dbl>, Petal.Length_max <dbl>,  
## #   Petal.Width_max <dbl>
```

More complex transformations

```
# You can express more complex inline transformations using .  
by_species %>% mutate_all(funs(. / 2.54))
```

```
## # A tibble: 150 x 5  
## # Groups:   Species [3]  
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species  
##           <dbl>         <dbl>         <dbl>         <dbl> <fct>  
## 1         2.01         1.38         0.551         0.0787 setosa  
## 2         1.93         1.18         0.551         0.0787 setosa  
## 3         1.85         1.26         0.512         0.0787 setosa  
## 4         1.81         1.22         0.591         0.0787 setosa  
## 5         1.97         1.42         0.551         0.0787 setosa  
## 6         2.13         1.54         0.669         0.157  setosa  
## 7         1.81         1.34         0.551         0.118  setosa  
## 8         1.97         1.34         0.591         0.0787 setosa  
## 9         1.73         1.14         0.551         0.0787 setosa  
## 10        1.93         1.22         0.591         0.0394 setosa  
## # ... with 140 more rows
```